

- How long does it take to complete the training run? (hint: this session is on distributed training, so it *will* take a while)
  - 19 hours and 36 mins
- Do you think your model is fully trained? How can you tell?
  - I think yes, because the training loss chart is almost a flat line
- Were you overfitting?
  - **Overfitting** if: training loss << eval loss, **Underfitting** if: training loss >> eval loss, **Just right** if eval loss ~ validation loss. To me both eval loss and training loss has the same shape, so we don't have overfitting.
- Were your GPUs fully utilized?
  - yes

```

X root@v100a: ~ (ssh)

```

	Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage		GPU-Util	Compute M.
0	N/A	38C	P0	43W / 250W	Off	00000000:00:07.0	Off	100%	Default
1	N/A	37C	P0	43W / 250W	Off	00000000:00:08.0	Off	100%	Default

```

X root@v100b: ~ (ssh)

```

GPU	Name	Persistence-MI	Bus-Id	Disp.A		Volatile	Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.
0	N/A	39C	P0	43W / 250W	Off	00000000:00:07.0	Off
1	N/A	39C	P0	42W / 250W	Off	00000000:00:08.0	Off


- Did you monitor network traffic (hint: `apt install nmon`)? Was network the bottleneck?
  - yes

root@v100a: ~ (ssh)									
nmmon-14g Hostname=v100a Refresh= 2secs 21:28.01									
Network I/O									
I/F	Name	Recv=KB/s	Trans=KB/s	packin	packout	insize	outsize	Peak->Recv	Trans
lo		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
eth1		0.1	0.3	2.5	1.0	47.2	358.0	5.7	14.2
docker0		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
eth0		236707.3	229114.5	168257.5	11461.1	1440.6	20470.4	1428723.2	1381516.8

root@v100a: ~ (ssh)									
nmmon-14g [H for help] Hostname=v100a Refresh= 2secs 02:57.27									
Network I/O									
I/F	Name	Recv=KB/s	Trans=KB/s	packin	packout	insize	outsize	Peak->Recv	Trans
lo		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
eth1		0.1	0.2	1.0	0.5	66.0	326.0	8.8	509.9
docker0		0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
eth0		248059.6	239955.3	176220.4	12099.0	1441.5	20308.6	1428723.2	1381516.8

- It is not, 1G bandwidth was not fully utilized.
- Take a look at the plot of the learning rate and then check the config file. Can you explain this setting?



- Learning rate is one of hyperparameters. Warm up steps are 8000, it allows the algorithm to go sharply from 0 to a high value before 8000 steps, then gradually decreases the learning rate to find the optimal value.
- How big was your training set (mb)? How many training lines did it contain?
- (I forgot to check for the qnswer before canceling the vs, recreated a p100 vs to answer this question!)

```

root@p100a:/data/wmt16_de_en# wc -l train.en
4562102 train.en
root@p100a:/data/wmt16_de_en# wc -l train.de
4562102 train.de
root@p100a:/data/wmt16_de_en# ls -l train.en
-rw-r--r-- 1 root root 636464546 Mar  6 03:20 train.en
root@p100a:/data/wmt16_de_en# ls -l train.de
-rw-r--r-- 1 root root 710264445 Mar  6 03:20 train.de
root@p100a:/data/wmt16_de_en# ls -sh train.en
607M train.en
root@p100a:/data/wmt16_de_en# ls -sh train.de
678M train.de

```

- What are the files that a TF checkpoint is comprised of?
  - each checkpoint has a .data, a .meta, and a .index
- How big is your resulting model checkpoint (mb)?
  - ~8M
- Remember the definition of a "step". How long did an average step take?
  - 2.5 sec (total steps were 4600 and run time was 19 hours and 36 minutes)
- How does that correlate with the observed network utilization between nodes?
  - the network utilization is positively correlated with the time required for each step, because with the increase time of communication between nodes, it takes more time to complete a step.